

Predict Stress Level By Sleeping Hours

Group 5 Member:

Vyos Devin (2501974776) Data Science, Binus University - vyos.devin@binus.ac.id (<mailto:vyos.devin@binus.ac.id>)

Lindawaty Veronica (2501972890) Data Science, Binus University - lindawaty.veronica@binus.ac.id (<mailto:lindawaty.veronica@binus.ac.id>)

Tifara Beata Wibowo (2501975223) Data Science, Binus University - tifara.wibowo@binus.ac.id (<mailto:tifara.wibowo@binus.ac.id>)

Vina Nathania Santiko (2501960985) Data Science, Binus University - vina.santiko@binus.ac.id (<mailto:vina.santiko@binus.ac.id>)

```
In [2]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split

from sklearn.tree import DecisionTreeClassifier
from sklearn.svm import SVC
from sklearn.neighbors import KNeighborsClassifier
import time
from sklearn.ensemble import RandomForestClassifier, VotingClassifier
from sklearn.model_selection import GridSearchCV

from sklearn.metrics import confusion_matrix, classification_report, accuracy_score

import warnings
warnings.filterwarnings("ignore")
```

```
In [55]: data = pd.read_csv("https://raw.githubusercontent.com/Lindaaa0/Dataset/master/SaYoPillow.csv")
data
```

Out[55]:

| | sr | rr | t | lm | bo | rem | sr.1 | hr | sl |
|-----|--------|--------|--------|--------|--------|--------|-------|-------|-----|
| 0 | 93.800 | 25.680 | 91.840 | 16.600 | 89.840 | 99.60 | 1.840 | 74.20 | 3 |
| 1 | 91.640 | 25.104 | 91.552 | 15.880 | 89.552 | 98.88 | 1.552 | 72.76 | 3 |
| 2 | 60.000 | 20.000 | 96.000 | 10.000 | 95.000 | 85.00 | 7.000 | 60.00 | 1 |
| 3 | 85.760 | 23.536 | 90.768 | 13.920 | 88.768 | 96.92 | 0.768 | 68.84 | 3 |
| 4 | 48.120 | 17.248 | 97.872 | 6.496 | 96.248 | 72.48 | 8.248 | 53.12 | 0 |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 625 | 69.600 | 20.960 | 92.960 | 10.960 | 90.960 | 89.80 | 3.440 | 62.40 | 2 |
| 626 | 48.440 | 17.376 | 98.064 | 6.752 | 96.376 | 73.76 | 8.376 | 53.44 | 0 |
| 627 | 97.504 | 27.504 | 86.880 | 17.752 | 84.256 | 101.88 | 0.000 | 78.76 | 4 |
| 628 | 58.640 | 19.728 | 95.728 | 9.728 | 94.592 | 84.32 | 6.728 | 59.32 | 1 |
| 629 | 73.920 | 21.392 | 93.392 | 11.392 | 91.392 | 91.96 | 4.088 | 63.48 | 2 |

630 rows × 9 columns

```
In [4]: data.columns = ['snoring_rate', 'respiration_rate', 'body_temperature', 'limb_movement',
                        'blood_oxygen', 'eye_movement', 'sleeping_hours', 'heart_rate', 'stress_level']
```

Exploratory Data Analysis (EDA)

```
In [5]: data.isnull().sum()
```

```
Out[5]: snoring_rate      0
respiration_rate      0
body_temperature      0
limb_movement         0
blood_oxygen          0
eye_movement          0
sleeping_hours        0
heart_rate            0
stress_level          0
dtype: int64
```

Tidak ada missing value pada tiap variabel/instances

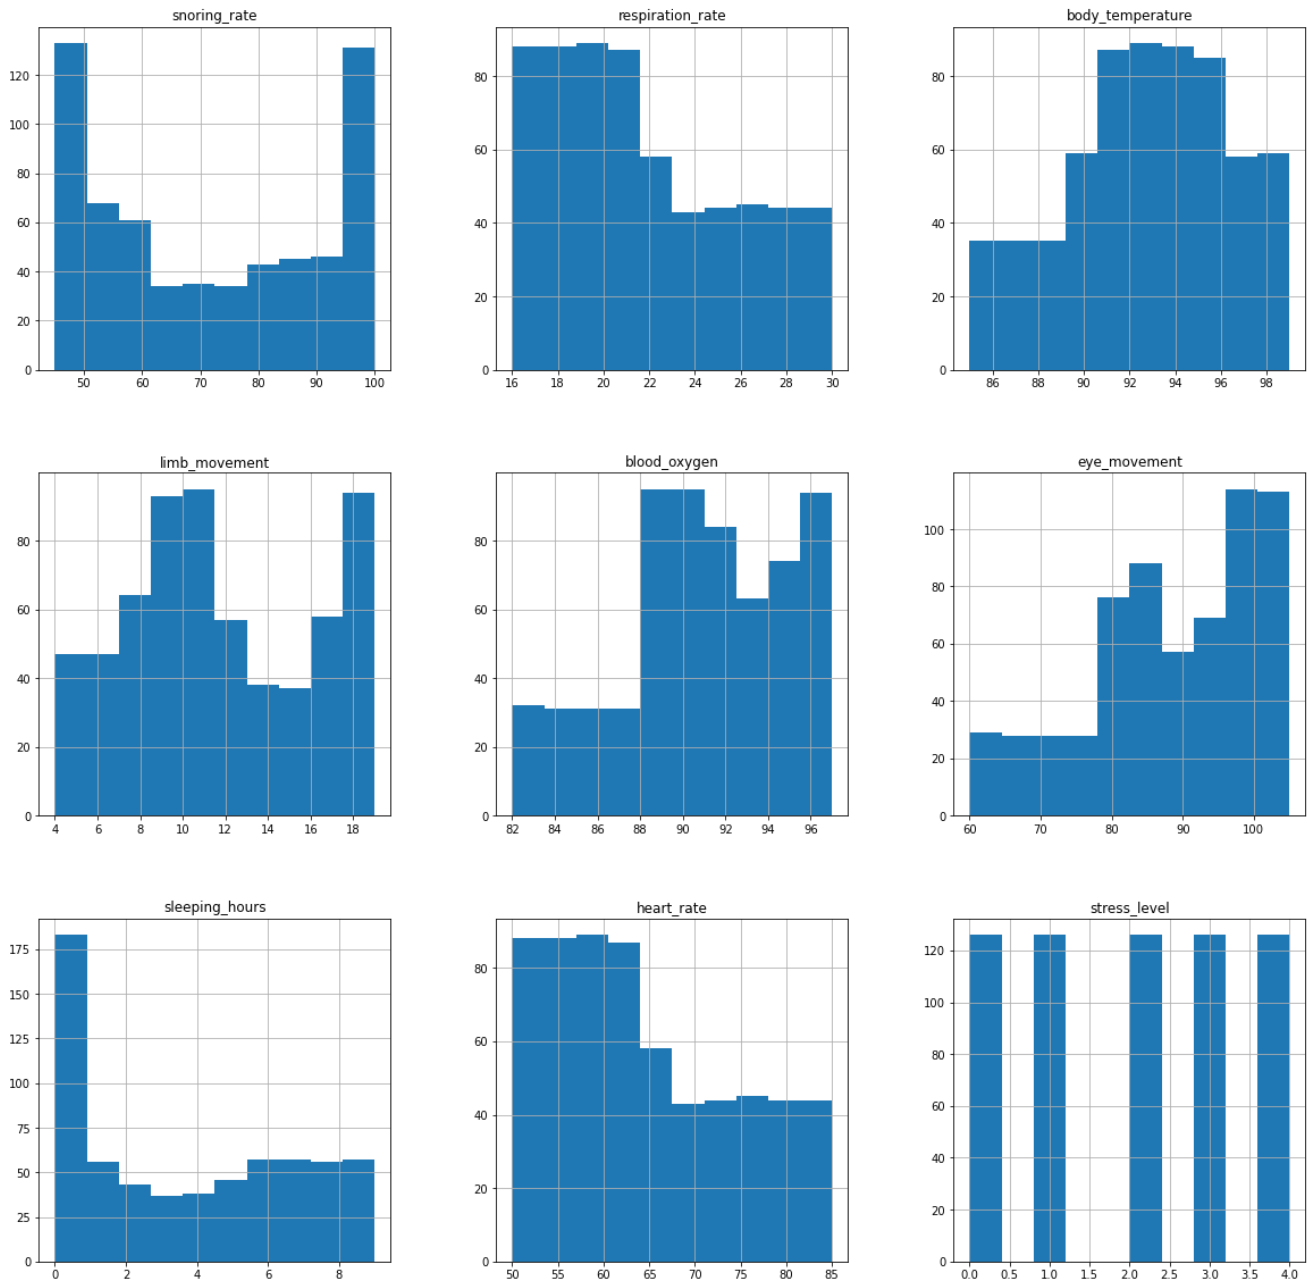
```
In [6]: data.describe()
```

Out[6]:

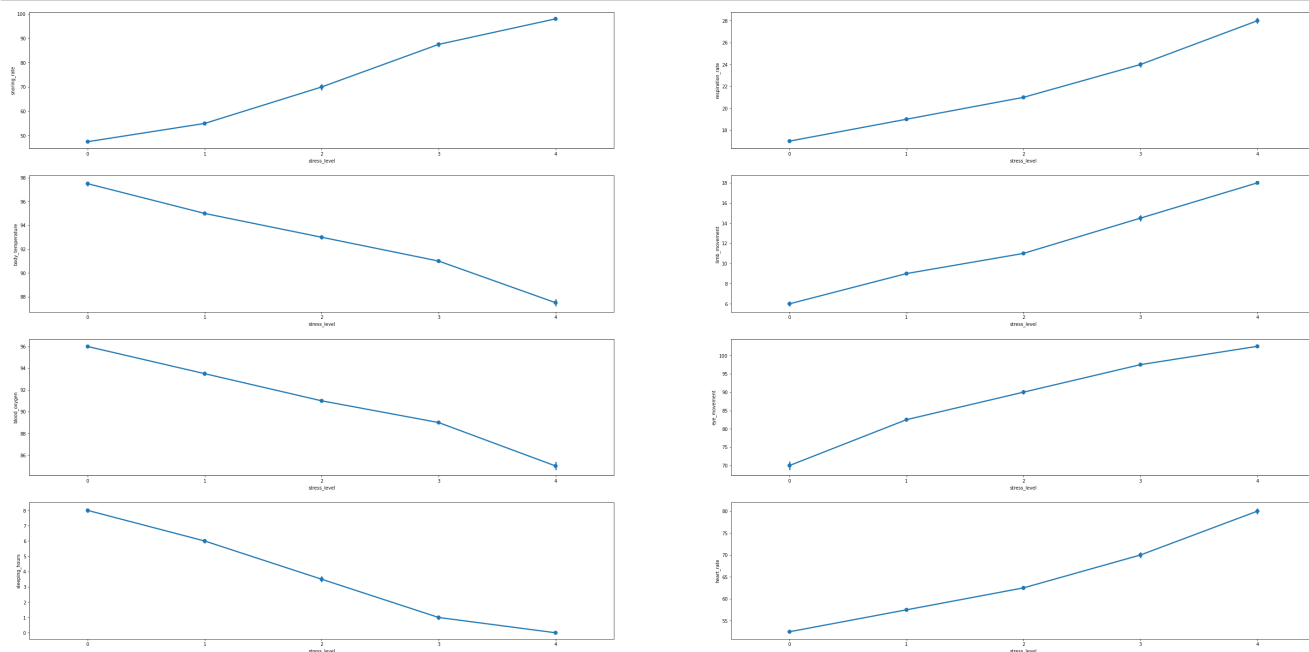
| | snoring_rate | respiration_rate | body_temperature | limb_movement | blood_oxygen | eye_movement | sleeping_hours | heart_rate | stress_level |
|-------|--------------|------------------|------------------|---------------|--------------|--------------|----------------|------------|--------------|
| count | 630.000000 | 630.000000 | 630.000000 | 630.000000 | 630.000000 | 630.000000 | 630.000000 | 630.000000 | 630.000000 |
| mean | 71.600000 | 21.800000 | 92.800000 | 11.700000 | 90.900000 | 88.500000 | 3.700000 | 64.500000 | 2.000000 |
| std | 19.372833 | 3.966111 | 3.52969 | 4.299629 | 3.902483 | 11.893747 | 3.054572 | 9.915277 | 1.415337 |
| min | 45.000000 | 16.000000 | 85.000000 | 4.000000 | 82.000000 | 60.000000 | 0.000000 | 50.000000 | 0.000000 |
| 25% | 52.500000 | 18.500000 | 90.500000 | 8.500000 | 88.500000 | 81.250000 | 0.500000 | 56.250000 | 1.000000 |
| 50% | 70.000000 | 21.000000 | 93.000000 | 11.000000 | 91.000000 | 90.000000 | 3.500000 | 62.500000 | 2.000000 |
| 75% | 91.250000 | 25.000000 | 95.500000 | 15.750000 | 94.250000 | 98.750000 | 6.500000 | 72.500000 | 3.000000 |
| max | 100.000000 | 30.000000 | 99.000000 | 19.000000 | 97.000000 | 105.000000 | 9.000000 | 85.000000 | 4.000000 |

Data Visualization

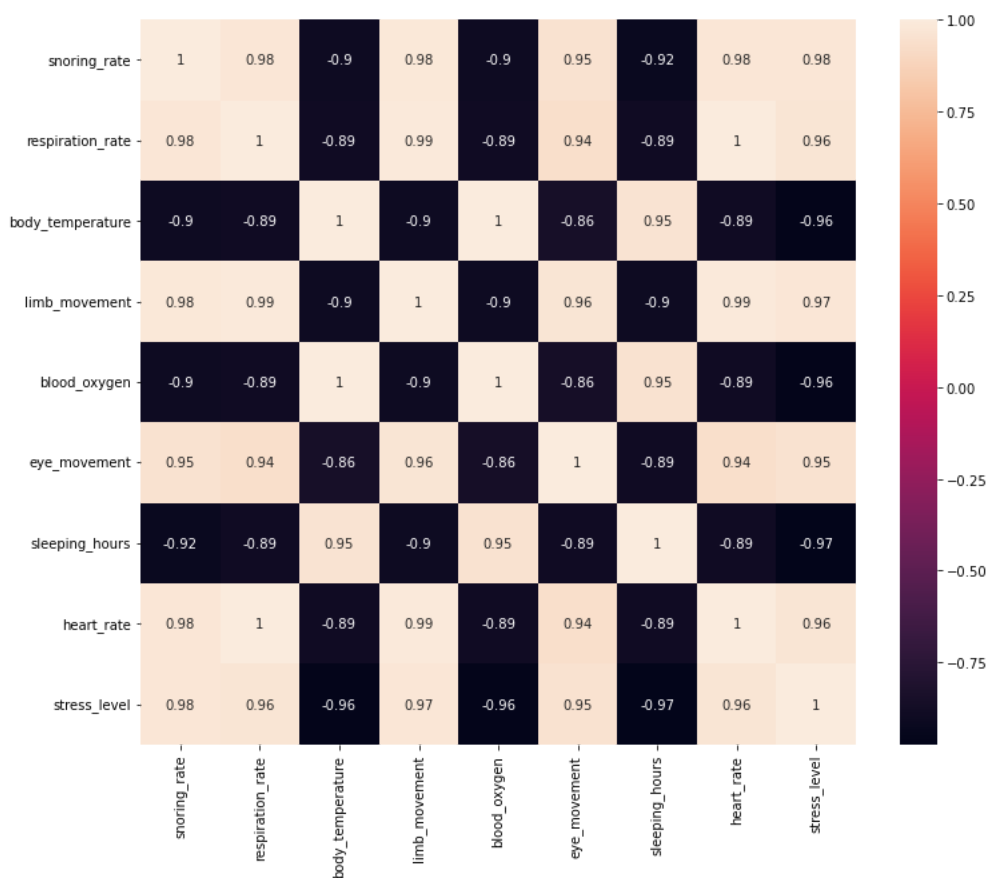
```
In [7]: dist = data.hist(figsize = (20, 20))
```



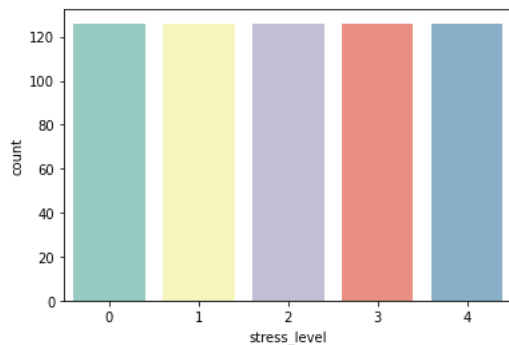
```
In [8]: fig = plt.figure(figsize=(50, 25))
rows = 4
columns = 2
for i in range(len(data.columns[:-1])):
    fig.add_subplot(rows, columns, (i+1))
    img = sns.pointplot(x='stress_level', y=data.columns[i], data=data)
plt.show()
```



```
In [9]: plt.figure(figsize = (12, 10))
p = sns.heatmap(data.corr(), annot = True)
```



```
In [10]: sns.countplot(x="stress_level", data = data, palette = "Set3")
plt.show()
```



Scaling Data

```
In [11]: #dropping dependent variable
sc = StandardScaler()
X = pd.DataFrame(sc.fit_transform(data.drop(["stress_level"],axis = 1)), columns=['snoring_rate', 'respiration_rate', 'body_temp',
                                          'blood_oxygen', 'eye_movement', 'sleeping_hours', 'heart_rate'])
X.head()
```

```
Out[11]:
```

| | snoring_rate | respiration_rate | body_temperature | limb_movement | blood_oxygen | eye_movement | sleeping_hours | heart_rate |
|---|--------------|------------------|------------------|---------------|--------------|--------------|----------------|------------|
| 0 | 1.146845 | 0.979066 | -0.272195 | 1.140539 | -0.271838 | 0.934005 | -0.609407 | 0.979066 |
| 1 | 1.035260 | 0.833720 | -0.353853 | 0.972949 | -0.345696 | 0.873421 | -0.703767 | 0.833720 |
| 2 | -0.599252 | -0.454206 | 0.907316 | -0.395697 | 1.051448 | -0.294506 | 1.081206 | -0.454206 |
| 3 | 0.731501 | 0.438056 | -0.576145 | 0.516734 | -0.546753 | 0.708498 | -0.960635 | 0.438056 |
| 4 | -1.212970 | -1.148636 | 1.438095 | -1.211299 | 1.371498 | -1.347997 | 1.490099 | -1.148636 |

```
In [12]: #assigning target variable
y = data.stress_level
y
```

```
Out[12]:
```

| | |
|-----|---|
| 0 | 3 |
| 1 | 3 |
| 2 | 1 |
| 3 | 3 |
| 4 | 0 |
| .. | |
| 625 | 2 |
| 626 | 0 |
| 627 | 4 |
| 628 | 1 |
| 629 | 2 |

Name: stress_level, Length: 630, dtype: int64

Model Building

```
In [31]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.2, random_state = 42)
```

Decision Tree

```
In [47]: dt = DecisionTreeClassifier()
dt.fit(X_train, y_train)
predict = dt.predict(X_test)
print("Confusion Matrix")
print("=====")
print(confusion_matrix(y_test, predict), "\n")

print("Classification Report")
print("=====")
print(classification_report(y_test, predict))
```

Confusion Matrix

```
=====
[[23  0  0  0  0]
 [ 1 23  0  0  0]
 [ 0  0 28  0  0]
 [ 0  0  0 25  1]
 [ 0  0  0  0 25]]
```

Classification Report

```
=====
              precision    recall  f1-score   support

     0       0.96         1.00         0.98         23
     1       1.00         0.96         0.98         24
     2       1.00         1.00         1.00         28
     3       1.00         0.96         0.98         26
     4       0.96         1.00         0.98         25

 accuracy          0.98
 macro avg         0.98         0.98         0.98
weighted avg         0.98         0.98         0.98
```


```
In [45]: train_scores, test_scores = list(), list()

values = [i for i in range(1, 21)]
for i in values:
    model = DecisionTreeClassifier(max_depth = i)
    model.fit(X_train, y_train)
    train_yhat = model.predict(X_train)
    train_acc = accuracy_score(y_train, train_yhat)
    train_scores.append(train_acc)
    test_yhat = model.predict(X_test)
    test_acc = accuracy_score(y_test, test_yhat)
    test_scores.append(test_acc)
    print('>%d, train: %.3f, test: %.3f' % (i, train_acc, test_acc))

# plot of train and test scores vs tree depth
plt.plot(values, train_scores, '-o', label='train_set', color = '#2B80FF')
plt.plot(values, test_scores, '-o', label='test_set', color = '#ff7e20')
plt.legend()
plt.title("The accuracy of train and test set\n", fontsize = 16)
plt.show()
```

```
>1, train: 0.407, test: 0.365
>2, train: 0.607, test: 0.556
>3, train: 0.806, test: 0.754
>4, train: 1.000, test: 0.976
>5, train: 1.000, test: 0.984
>6, train: 1.000, test: 0.984
>7, train: 1.000, test: 0.984
>8, train: 1.000, test: 0.984
>9, train: 1.000, test: 0.984
>10, train: 1.000, test: 0.984
>11, train: 1.000, test: 0.976
>12, train: 1.000, test: 0.976
>13, train: 1.000, test: 0.984
>14, train: 1.000, test: 0.984
>15, train: 1.000, test: 0.984
>16, train: 1.000, test: 0.984
>17, train: 1.000, test: 0.984
>18, train: 1.000, test: 0.984
>19, train: 1.000, test: 0.984
>20, train: 1.000, test: 0.984
```

The accuracy of train and test set



SVM

```
In [16]: svc_model = SVC()
svc_model.fit(X_train, y_train)

svc_pred = svc_model.predict(X_test)
print(confusion_matrix(y_test, svc_pred))
print(classification_report(y_test, svc_pred))
```

```
[[16  0  0  0  0]
 [ 0 10  0  0  0]
 [ 0  0 14  0  0]
 [ 0  0  0 14  0]
 [ 0  0  0  0  9]]
```


| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 1.00 | 1.00 | 1.00 | 16 |
| 1 | 1.00 | 1.00 | 1.00 | 10 |
| 2 | 1.00 | 1.00 | 1.00 | 14 |
| 3 | 1.00 | 1.00 | 1.00 | 14 |
| 4 | 1.00 | 1.00 | 1.00 | 9 |
| accuracy | | | 1.00 | 63 |
| macro avg | 1.00 | 1.00 | 1.00 | 63 |
| weighted avg | 1.00 | 1.00 | 1.00 | 63 |

```
In [17]: train_scores, test_scores = list(), list()

values = [i for i in range(1, 21)]
for i in values:
    model = SVC()
    model.fit(X_train, y_train)
    train_yhat = model.predict(X_train)
    train_acc = accuracy_score(y_train, train_yhat)
    train_scores.append(train_acc)
    test_yhat = model.predict(X_test)
    test_acc = accuracy_score(y_test, test_yhat)
    test_scores.append(test_acc)
    print('>%d, train: %.3f, test: %.3f' % (i, train_acc, test_acc))

# plot of train and test scores vs tree depth
plt.plot(values, train_scores, '-o', label='train_set')
plt.plot(values, test_scores, '-o', label='test_set')
plt.legend()
plt.show()
```

```
>1, train: 1.000, test: 1.000
>2, train: 1.000, test: 1.000
>3, train: 1.000, test: 1.000
>4, train: 1.000, test: 1.000
>5, train: 1.000, test: 1.000
>6, train: 1.000, test: 1.000
>7, train: 1.000, test: 1.000
>8, train: 1.000, test: 1.000
>9, train: 1.000, test: 1.000
>10, train: 1.000, test: 1.000
>11, train: 1.000, test: 1.000
>12, train: 1.000, test: 1.000
>13, train: 1.000, test: 1.000
>14, train: 1.000, test: 1.000
>15, train: 1.000, test: 1.000
>16, train: 1.000, test: 1.000
>17, train: 1.000, test: 1.000
>18, train: 1.000, test: 1.000
>19, train: 1.000, test: 1.000
>20, train: 1.000, test: 1.000
```



KNN

```
In [18]: knn = KNeighborsClassifier()
knnModel = knn.fit(X_train, y_train)
knnPred = knn.predict(X_train)
print(confusion_matrix(y_train, knnPred))
print(classification_report(y_train, knnPred))
```


```
knnPred = knn.predict(X_test)
print(confusion_matrix(y_test, knnPred))
print(classification_report(y_test, knnPred))
```

| | | | | | |
|--------------------|-----------|--------|----------|---------|--|
| [[110 0 0 0 0] | | | | | |
| [0 116 0 0 0] | | | | | |
| [0 0 112 0 0] | | | | | |
| [0 0 0 112 0] | | | | | |
| [0 0 0 0 117]] | | | | | |
| | precision | recall | f1-score | support | |
| 0 | 1.00 | 1.00 | 1.00 | 110 | |
| 1 | 1.00 | 1.00 | 1.00 | 116 | |
| 2 | 1.00 | 1.00 | 1.00 | 112 | |
| 3 | 1.00 | 1.00 | 1.00 | 112 | |
| 4 | 1.00 | 1.00 | 1.00 | 117 | |
| accuracy | | | 1.00 | 567 | |
| macro avg | 1.00 | 1.00 | 1.00 | 567 | |
| weighted avg | 1.00 | 1.00 | 1.00 | 567 | |

| | | | | | |
|-------------------|-----------|--------|----------|---------|--|
| [[16 0 0 0 0] | | | | | |
| [0 10 0 0 0] | | | | | |
| [0 0 14 0 0] | | | | | |
| [0 0 0 14 0] | | | | | |
| [0 0 0 0 9]] | | | | | |
| | precision | recall | f1-score | support | |
| 0 | 1.00 | 1.00 | 1.00 | 16 | |
| 1 | 1.00 | 1.00 | 1.00 | 10 | |
| 2 | 1.00 | 1.00 | 1.00 | 14 | |
| 3 | 1.00 | 1.00 | 1.00 | 14 | |
| 4 | 1.00 | 1.00 | 1.00 | 9 | |
| accuracy | | | 1.00 | 63 | |
| macro avg | 1.00 | 1.00 | 1.00 | 63 | |
| weighted avg | 1.00 | 1.00 | 1.00 | 63 | |


```
In [19]: train_scores, test_scores = list(), list()
values = [i for i in range(1, 51)]
for i in values:
    model = KNeighborsClassifier(n_neighbors=i)
    model.fit(X_train, y_train)
    train_yhat = model.predict(X_train)
    train_acc = accuracy_score(y_train, train_yhat)
    train_scores.append(train_acc)
    test_yhat = model.predict(X_test)
    test_acc = accuracy_score(y_test, test_yhat)
    test_scores.append(test_acc)
    print('>%d, train: %.3f, test: %.3f' % (i, train_acc, test_acc))
# plot of train and test scores vs number of neighbors
plt.plot(values, train_scores, '-o', label='Train')
plt.plot(values, test_scores, '-o', label='Test')
plt.legend()
plt.show()
```

```
>1, train: 1.000, test: 1.000
>2, train: 1.000, test: 1.000
>3, train: 1.000, test: 1.000
>4, train: 1.000, test: 1.000
>5, train: 1.000, test: 1.000
>6, train: 1.000, test: 1.000
>7, train: 1.000, test: 1.000
>8, train: 1.000, test: 1.000
>9, train: 1.000, test: 1.000
>10, train: 1.000, test: 1.000
>11, train: 1.000, test: 1.000
>12, train: 1.000, test: 1.000
>13, train: 1.000, test: 1.000
>14, train: 1.000, test: 1.000
>15, train: 1.000, test: 1.000
>16, train: 1.000, test: 1.000
>17, train: 1.000, test: 1.000
>18, train: 1.000, test: 1.000
>19, train: 1.000, test: 1.000
>20, train: 1.000, test: 1.000
>21, train: 1.000, test: 1.000
>22, train: 1.000, test: 1.000
>23, train: 1.000, test: 1.000
>24, train: 1.000, test: 1.000
>25, train: 1.000, test: 1.000
>26, train: 1.000, test: 1.000
>27, train: 1.000, test: 1.000
>28, train: 1.000, test: 1.000
>29, train: 1.000, test: 1.000
>30, train: 1.000, test: 1.000
>31, train: 1.000, test: 1.000
>32, train: 1.000, test: 1.000
>33, train: 1.000, test: 1.000
>34, train: 1.000, test: 1.000
>35, train: 1.000, test: 1.000
>36, train: 1.000, test: 1.000
>37, train: 1.000, test: 1.000
>38, train: 1.000, test: 1.000
>39, train: 1.000, test: 1.000
>40, train: 1.000, test: 1.000
>41, train: 1.000, test: 1.000
>42, train: 1.000, test: 1.000
>43, train: 1.000, test: 1.000
>44, train: 1.000, test: 1.000
>45, train: 1.000, test: 1.000
>46, train: 1.000, test: 1.000
>47, train: 1.000, test: 1.000
>48, train: 1.000, test: 1.000
>49, train: 1.000, test: 1.000
>50, train: 1.000, test: 1.000
```



Voting Classifier

```
In [22]: estimators = [('svm', svc_model), ('KNN', knn), ('DecisionTree', dt)]

ensemble = VotingClassifier(estimators, voting = 'hard')


ensembleModel = ensemble.fit(X_train, y_train)

ensemblePred = ensemble.predict(X_test)
print(classification_report(y_test, ensemblePred))
```

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 1.00 | 1.00 | 1.00 | 16 |
| 1 | 1.00 | 1.00 | 1.00 | 10 |
| 2 | 1.00 | 1.00 | 1.00 | 14 |
| 3 | 1.00 | 1.00 | 1.00 | 14 |
| 4 | 1.00 | 1.00 | 1.00 | 9 |
| accuracy | | | 1.00 | 63 |
| macro avg | 1.00 | 1.00 | 1.00 | 63 |
| weighted avg | 1.00 | 1.00 | 1.00 | 63 |

```
In [23]: train_scores, test_scores = list(), list()
values = [i for i in range(1, 51)]
for i in values:
    model = ensemble
    model.fit(X_train, y_train)
    train_yhat = model.predict(X_train)
    train_acc = accuracy_score(y_train, train_yhat)
    train_scores.append(train_acc)
    test_yhat = model.predict(X_test)
    test_acc = accuracy_score(y_test, test_yhat)
    test_scores.append(test_acc)
    print('>%d, train: %.3f, test: %.3f' % (i, train_acc, test_acc))
# plot of train and test scores vs number of neighbors
plt.plot(values, train_scores, '-o', label='Train')
plt.plot(values, test_scores, '-o', label='Test')
plt.legend()
plt.show()
```

```
>1, train: 1.000, test: 1.000
>2, train: 1.000, test: 1.000
>3, train: 1.000, test: 1.000
>4, train: 1.000, test: 1.000
>5, train: 1.000, test: 1.000
>6, train: 1.000, test: 1.000
>7, train: 1.000, test: 1.000
>8, train: 1.000, test: 1.000
>9, train: 1.000, test: 1.000
>10, train: 1.000, test: 1.000
>11, train: 1.000, test: 1.000
>12, train: 1.000, test: 1.000
>13, train: 1.000, test: 1.000
>14, train: 1.000, test: 1.000
>15, train: 1.000, test: 1.000
>16, train: 1.000, test: 1.000
>17, train: 1.000, test: 1.000
>18, train: 1.000, test: 1.000
>19, train: 1.000, test: 1.000
>20, train: 1.000, test: 1.000
>21, train: 1.000, test: 1.000
>22, train: 1.000, test: 1.000
>23, train: 1.000, test: 1.000
>24, train: 1.000, test: 1.000
>25, train: 1.000, test: 1.000
>26, train: 1.000, test: 1.000
>27, train: 1.000, test: 1.000
>28, train: 1.000, test: 1.000
>29, train: 1.000, test: 1.000
>30, train: 1.000, test: 1.000
>31, train: 1.000, test: 1.000
>32, train: 1.000, test: 1.000
>33, train: 1.000, test: 1.000
>34, train: 1.000, test: 1.000
>35, train: 1.000, test: 1.000
>36, train: 1.000, test: 1.000
>37, train: 1.000, test: 1.000
>38, train: 1.000, test: 1.000
>39, train: 1.000, test: 1.000
>40, train: 1.000, test: 1.000
>41, train: 1.000, test: 1.000
>42, train: 1.000, test: 1.000
>43, train: 1.000, test: 1.000
>44, train: 1.000, test: 1.000
>45, train: 1.000, test: 1.000
>46, train: 1.000, test: 1.000
>47, train: 1.000, test: 1.000
>48, train: 1.000, test: 1.000
>49, train: 1.000, test: 1.000
>50, train: 1.000, test: 1.000
```



Bagging

```
In [20]: start = time.time()
param_dist = {'max_depth' : [2,3,4],
              'bootstrap': [True, False],
              'max_features': ['auto', 'sqrt', 'log2', None],
              'criterion': ['gini', 'entropy']}

fit_rf = RandomForestClassifier()
cv_rf = GridSearchCV(fit_rf, cv = 10, param_grid = param_dist, n_jobs = 3)

rfModel = cv_rf.fit(X_train, y_train)
print('Best Parameters using grid search: \n', cv_rf.best_params_)
end = time.time()
print('Time taken in grid search: {0: .2f}'.format(end - start))
```

Best Parameters using grid search:
 {'bootstrap': True, 'criterion': 'entropy', 'max_depth': 2, 'max_features': 'auto'}
 Time taken in grid search: 61.92

```
In [21]: rfPred = cv_rf.predict(X_test)
print(classification_report(y_test, rfPred))
```

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 1.00 | 1.00 | 1.00 | 16 |
| 1 | 1.00 | 0.90 | 0.95 | 10 |
| 2 | 0.93 | 1.00 | 0.97 | 14 |
| 3 | 1.00 | 0.93 | 0.96 | 14 |
| 4 | 0.90 | 1.00 | 0.95 | 9 |
| accuracy | | | 0.97 | 63 |
| macro avg | 0.97 | 0.97 | 0.96 | 63 |
| weighted avg | 0.97 | 0.97 | 0.97 | 63 |

```
In [27]: rfPred = cv_rf.predict(X_train)
print(classification_report(y_train, rfPred))

rfPred = cv_rf.predict(X_test)
print(classification_report(y_test, rfPred))
```

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 1.00 | 1.00 | 1.00 | 110 |
| 1 | 1.00 | 1.00 | 1.00 | 116 |
| 2 | 1.00 | 1.00 | 1.00 | 112 |
| 3 | 1.00 | 1.00 | 1.00 | 112 |
| 4 | 1.00 | 1.00 | 1.00 | 117 |
| accuracy | | | 1.00 | 567 |
| macro avg | 1.00 | 1.00 | 1.00 | 567 |
| weighted avg | 1.00 | 1.00 | 1.00 | 567 |

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.94 | 1.00 | 0.97 | 16 |
| 1 | 1.00 | 0.90 | 0.95 | 10 |
| 2 | 1.00 | 1.00 | 1.00 | 14 |
| 3 | 1.00 | 0.93 | 0.96 | 14 |
| 4 | 0.90 | 1.00 | 0.95 | 9 |
| accuracy | | | 0.97 | 63 |
| macro avg | 0.97 | 0.97 | 0.97 | 63 |
| weighted avg | 0.97 | 0.97 | 0.97 | 63 |

```

In [25]: train_scores, test_scores = list(), list()
values = [i for i in range(1, 21)]
for i in values:
    model = cv_rf
    model.fit(X_train, y_train)
    train_yhat = model.predict(X_train)
    train_acc = accuracy_score(y_train, train_yhat)
    train_scores.append(train_acc)
    test_yhat = model.predict(X_test)
    test_acc = accuracy_score(y_test, test_yhat)
    test_scores.append(test_acc)
    print('>%d, train: %.3f, test: %.3f' % (i, train_acc, test_acc))
# plot of train and test scores vs number of neighbors
plt.plot(values, train_scores, '-o', label='Train')
plt.plot(values, test_scores, '-o', label='Test')
plt.legend()
plt.show()

```

```

>1, train: 1.000, test: 0.968
>2, train: 1.000, test: 0.968
>3, train: 1.000, test: 0.968
>4, train: 1.000, test: 0.968
>5, train: 1.000, test: 0.968
>6, train: 1.000, test: 0.968
>7, train: 1.000, test: 0.968
>8, train: 1.000, test: 0.968
>9, train: 1.000, test: 0.968
>10, train: 1.000, test: 0.968
>11, train: 1.000, test: 0.968
>12, train: 1.000, test: 0.968
>13, train: 1.000, test: 0.968
>14, train: 1.000, test: 0.968
>15, train: 1.000, test: 0.968
>16, train: 1.000, test: 0.968
>17, train: 1.000, test: 0.968
>18, train: 1.000, test: 0.968
>19, train: 1.000, test: 0.968
>20, train: 1.000, test: 0.968

```

